

a-task1-titanic-classification

August 12, 2024

CodeAlpha - Task1 - Titanic Classification

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix
```

```
[2]: # Load the dataset
ds = pd.read_csv("titanic.csv")

# Display the first few rows of the dataset
print(ds.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S

3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

Data Preprocessing

```
[6]: ds.columns
```

```
[6]: Index(['PassengerId', 'Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch',
          'Fare', 'Embarked'],
          dtype='object')
```

```
[8]: for i in ds.columns:
      if ds[f"{i}"].sum()==0:
          ds.drop(f"{i}", axis=1, inplace=True)
      ds.head()
```

```
[8]:
```

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	1	0	3	0	22.0	1	0	7.2500	0
1	2	1	1	1	38.0	1	0	71.2833	1
2	3	1	3	1	26.0	0	0	7.9250	0
3	4	1	1	1	35.0	1	0	53.1000	0
4	5	0	3	0	35.0	0	0	8.0500	0

```
[9]: ds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Sex             891 non-null   int64
4   Age             891 non-null   float64
5   SibSp           891 non-null   int64
6   Parch           891 non-null   int64
7   Fare            891 non-null   float64
8   Embarked        891 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 62.8 KB
```

```
[10]: ds.describe()
```

```
[10]:
```

	PassengerId	Survived	Pclass	Sex	Age	\
count	891.000000	891.000000	891.000000	891.000000	891.000000	
mean	446.000000	0.383838	2.308642	0.352413	29.361582	
std	257.353842	0.486592	0.836071	0.477990	13.019697	

min	1.000000	0.000000	1.000000	0.000000	0.420000
25%	223.500000	0.000000	2.000000	0.000000	22.000000
50%	446.000000	0.000000	3.000000	0.000000	28.000000
75%	668.500000	1.000000	3.000000	1.000000	35.000000
max	891.000000	1.000000	3.000000	1.000000	80.000000

	SibSp	Parch	Fare	Embarked
count	891.000000	891.000000	891.000000	891.000000
mean	0.523008	0.381594	32.204208	0.361392
std	1.102743	0.806057	49.693429	0.635673
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	7.910400	0.000000
50%	0.000000	0.000000	14.454200	0.000000
75%	1.000000	0.000000	31.000000	1.000000
max	8.000000	6.000000	512.329200	2.000000


```
[11]: ds.isnull().sum()
```

```
[11]: PassengerId    0
      Survived      0
      Pclass        0
      Sex           0
      Age           0
      SibSp         0
      Parch         0
      Fare          0
      Embarked      0
      dtype: int64
```

```
[12]: ds.fillna(ds.Embarked.mean(), inplace = True)
      ds.isnull().sum()
```

```
[12]: PassengerId    0
      Survived      0
      Pclass        0
      Sex           0
      Age           0
      SibSp         0
      Parch         0
      Fare          0
      Embarked      0
      dtype: int64
```

```
[13]: plt.figure(figsize=(20, 15))
      sns.heatmap(ds.corr(), annot=True, cmap='Blues')
      plt.title("Corelation", size=30)
      plt.show()
```



```
[14]: ds.head()
```

```
[14]:   PassengerId  Survived  Pclass  Sex  Age  SibSp  Parch    Fare  Embarked
0           1         0       3    0  22.0     1     0   7.2500         0
1           2         1       1    1  38.0     1     0  71.2833         1
2           3         1       3    1  26.0     0     0   7.9250         0
3           4         1       1    1  35.0     1     0  53.1000         0
4           5         0       3    0  35.0     0     0   8.0500         0
```

Logistic Regression

```
[22]: # Logistic Regression
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred_logreg = logreg.predict(X_test)

# Evaluation
print("Logistic Regression")
```

```

print("Accuracy:", accuracy_score(y_test, y_pred_logreg))
print(classification_report(y_test, y_pred_logreg))


# Confusion Matrix
sns.heatmap(confusion_matrix(y_test, y_pred_logreg), annot=True, fmt='d',
            cmap='Blues')
plt.title('Confusion Matrix - Logistic Regression')
plt.show()

```

Logistic Regression

Accuracy: 0.7988826815642458

	precision	recall	f1-score	support
0	0.82	0.85	0.83	105
1	0.77	0.73	0.75	74
accuracy			0.80	179
macro avg	0.79	0.79	0.79	179
weighted avg	0.80	0.80	0.80	179



Decision Tree

```
[20]: # Decision Tree
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)
y_pred_dtree = dtree.predict(X_test)

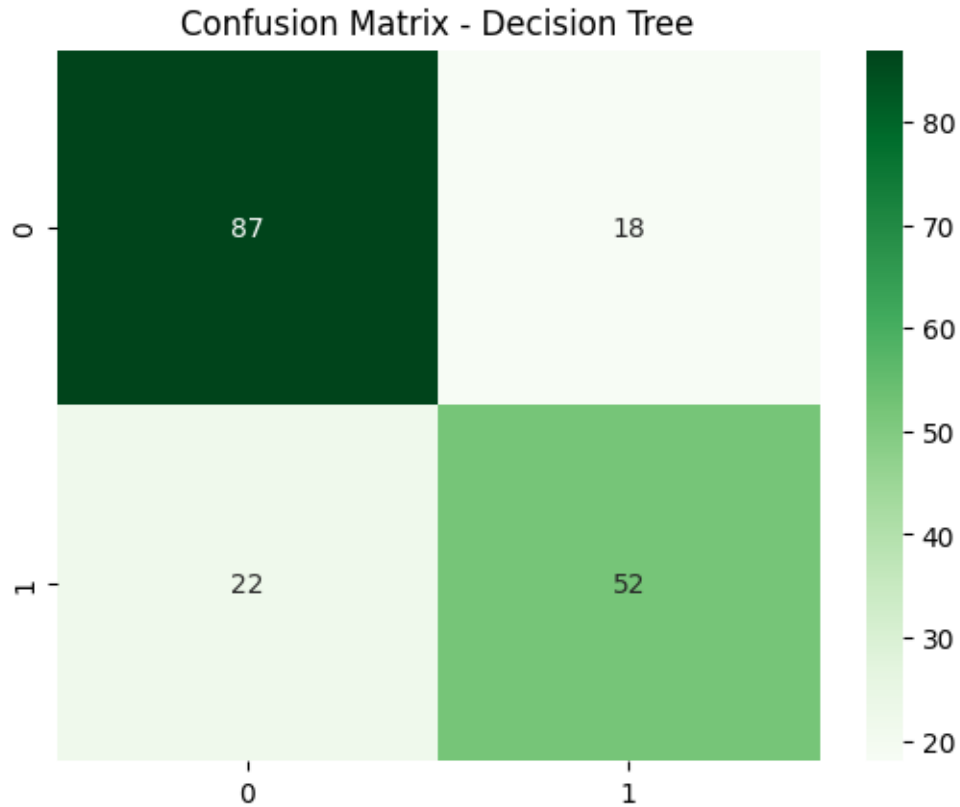
# Evaluation
print("Decision Tree")
print("Accuracy:", accuracy_score(y_test, y_pred_dtree))
print(classification_report(y_test, y_pred_dtree))

# Confusion Matrix
sns.heatmap(confusion_matrix(y_test, y_pred_dtree), annot=True, fmt='d', cmap='Greens')
plt.title('Confusion Matrix - Decision Tree')
plt.show()
```

Decision Tree

Accuracy: 0.776536312849162

	precision	recall	f1-score	support
0	0.80	0.83	0.81	105
1	0.74	0.70	0.72	74
accuracy			0.78	179
macro avg	0.77	0.77	0.77	179
weighted avg	0.78	0.78	0.78	179



Random Forest

```
[23]: # Random Forest
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

# Evaluation
print("Random Forest")
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))


# Confusion Matrix
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, fmt='d',
            cmap='Oranges')
plt.title('Confusion Matrix - Random Forest')
plt.show()
```

Random Forest

Accuracy: 0.8156424581005587

	precision	recall	f1-score	support
--	-----------	--------	----------	---------


0	0.82	0.88	0.85	105
1	0.81	0.73	0.77	74
accuracy			0.82	179
macro avg	0.81	0.80	0.81	179
weighted avg	0.82	0.82	0.81	179



Visualization

Pie Chart - Survival Rate


```
[25]: # Pie chart of survival rate
survived_counts = ds['Survived'].value_counts()
plt.pie(survived_counts, labels=['Not Survived', 'Survived'], autopct='%1.1f%%', colors=['lightcoral', 'skyblue'])
plt.title('Survival Rate')
plt.show()
```

Donut Chart - Gender Distribution


```
[28]: # Donut chart of gender distribution
gender_counts = ds['Sex'].value_counts()
plt.pie(gender_counts, labels=['Male', 'Female'], autopct='%1.1f%%',
        colors=['lightblue', 'Red'], wedgeprops=dict(width=0.3))
plt.title('Gender Distribution')
plt.show()
```

Gender Distribution



Scatter Plot - Age vs. Fare

```
[30]: # Scatter plot of Age vs Fare
plt.scatter(ds['Age'], ds['Fare'], c=ds['Survived'], cmap='coolwarm', alpha=0.9)
plt.title('Age vs Fare')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.colorbar(label='Survived')
plt.show()
```



```
[34]: # Sample DataFrame (assuming `ds` is already defined)
# Check column names and initial rows to ensure they are correct
print("Column Names:")
print(ds.columns)
print("\nSample Data:")
print(ds.head())
```

Column Names:

```
Index(['PassengerId', 'Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch',
      'Fare', 'Embarked'],
      dtype='object')
```

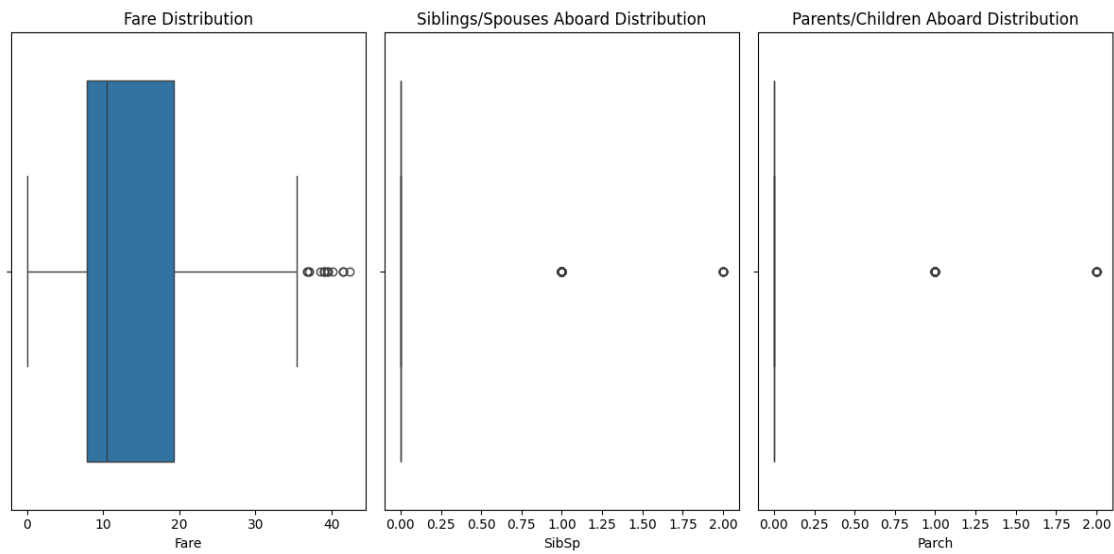
Sample Data:

	PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked
0	1	0	3	0	22.0	1	0	7.2500	0
2	3	1	3	1	26.0	0	0	7.9250	0
4	5	0	3	0	35.0	0	0	8.0500	0
5	6	0	3	0	28.0	0	0	8.4583	2
7	8	0	3	0	2.0	3	1	21.0750	0


```
[37]: ds.columns
```

```
[37]: Index(['PassengerId', 'Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch',  
        'Fare', 'Embarked'],  
        dtype='object')
```

```
[38]: # Dropping outliers  
ds = ds[ds['Fare'] < 45]  
ds = ds[ds['SibSp'] < 3] # Changed 'sibsp' to 'SibSp'  
ds = ds[ds['Parch'] < 3]  
  
# Plot boxplots for numerical features  
plt.figure(figsize=(12, 6))  
  
# Boxplot for 'Fare'  
plt.subplot(1, 3, 1)  
sns.boxplot(x=ds['Fare'])  
plt.title('Fare Distribution')  
  
# Boxplot for 'SibSp'  
plt.subplot(1, 3, 2)  
sns.boxplot(x=ds['SibSp']) # Changed 'sibsp' to 'SibSp'  
plt.title('Siblings/Spouses Aboard Distribution')  
  
# Boxplot for 'Parch'  
plt.subplot(1, 3, 3)  
sns.boxplot(x=ds['Parch'])  
plt.title('Parents/Children Aboard Distribution')  
  
plt.tight_layout()  
plt.show()  
  
# Pie chart for target variable distribution  
labels = ['Not Survived', 'Survived']  
sizes = ds['Survived'].value_counts() # Assuming 'Survived' should be  
    ↳ 'Survived'  
colors = ['#ff9999', '#66b3ff']  
plt.pie(sizes, labels=labels, colors=colors, autopct='%1.1f%%', startangle=140)  
plt.title('Survival Distribution')  
plt.show()
```



Survival Distribution



```
[39]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler

      # Assuming 'Survived' is the target variable and others are features
```

```

X = ds.drop('Survived', axis=1)
y = ds['Survived']

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Standardize features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

```

[40]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, classification_report

      # Initialize the model
      model = LogisticRegression()

      # Train the model
      model.fit(X_train, y_train)

      # Predict on the test set
      y_pred = model.predict(X_test)

      # Evaluate the model
      print("Accuracy:", accuracy_score(y_test, y_pred))
      print("Classification Report:\n", classification_report(y_test, y_pred))

```

Accuracy: 0.8602941176470589

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.92	0.90	90
1	0.83	0.74	0.78	46
accuracy			0.86	136
macro avg	0.85	0.83	0.84	136
weighted avg	0.86	0.86	0.86	136

```

[41]: from sklearn.tree import DecisionTreeClassifier

      # Initialize the model
      model = DecisionTreeClassifier()

      # Train the model
      model.fit(X_train, y_train)

```

```

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

```

Accuracy: 0.7352941176470589

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.81	0.80	90
1	0.61	0.59	0.60	46
accuracy			0.74	136
macro avg	0.70	0.70	0.70	136
weighted avg	0.73	0.74	0.73	136

```

[42]: from sklearn.ensemble import RandomForestClassifier

# Initialize the model
model = RandomForestClassifier()

# Train the model
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))

```

Accuracy: 0.8676470588235294

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.97	0.91	90
1	0.91	0.67	0.78	46
accuracy			0.87	136
macro avg	0.88	0.82	0.84	136
weighted avg	0.87	0.87	0.86	136

[]: